

PyTorch Tutorial

-NTU Machine Learning Course-

Lyman Lin 林裕訓

Nov. 03, 2017

lymanblue[at]gmail.com

What is PyTorch?

- Developed by Facebook
 - Python first
 - Dynamic Neural Network
 - This tutorial is for PyTorch 0.2.0
- Endorsed by Director of AI at Tesla



Andrej Karpathy 
@karpathy

Follow



I've been using PyTorch a few months now and I've never felt better. I have more energy. My skin is clearer. My eye sight has improved.

Installation

- PyTorch Web: <http://pytorch.org/>

Get Started.

Select your preferences, then run the PyTorch install command.

Please ensure that you are on the latest pip and numpy packages.
Anaconda is our recommended package manager

| | | | |
|-----------------|--------------|------------|--------|
| OS | Linux | OSX | |
| Package Manager | conda | pip | Source |
| Python | 2.7 | 3.5 | 3.6 |
| CUDA | 7.5 | 8.0 | None |

Run this command:

```
conda install pytorch torchvision cuda80 -c soumith
```

Packages of PyTorch

This Tutorial

| Package | Description |
|--------------------------|--|
| torch | a Tensor library like Numpy, with strong GPU support |
| torch.autograd | a tape based automatic differentiation library that supports all differentiable Tensor operations in torch |
| torch.nn | a neural networks library deeply integrated with autograd designed for maximum flexibility |
| torch.optim | an optimization package to be used with torch.nn with standard optimization methods such as SGD, RMSProp, LBFGS, Adam etc. |
| torch.multiprocessing | python multiprocessing, but with magical memory sharing of torch Tensors across processes. Useful for data loading and hogwild training. |
| torch.utils | DataLoader, Trainer and other utility functions for convenience |
| torch.legacy(.nn/.optim) | legacy code that has been ported over from torch for backward compatibility reasons |

Outline

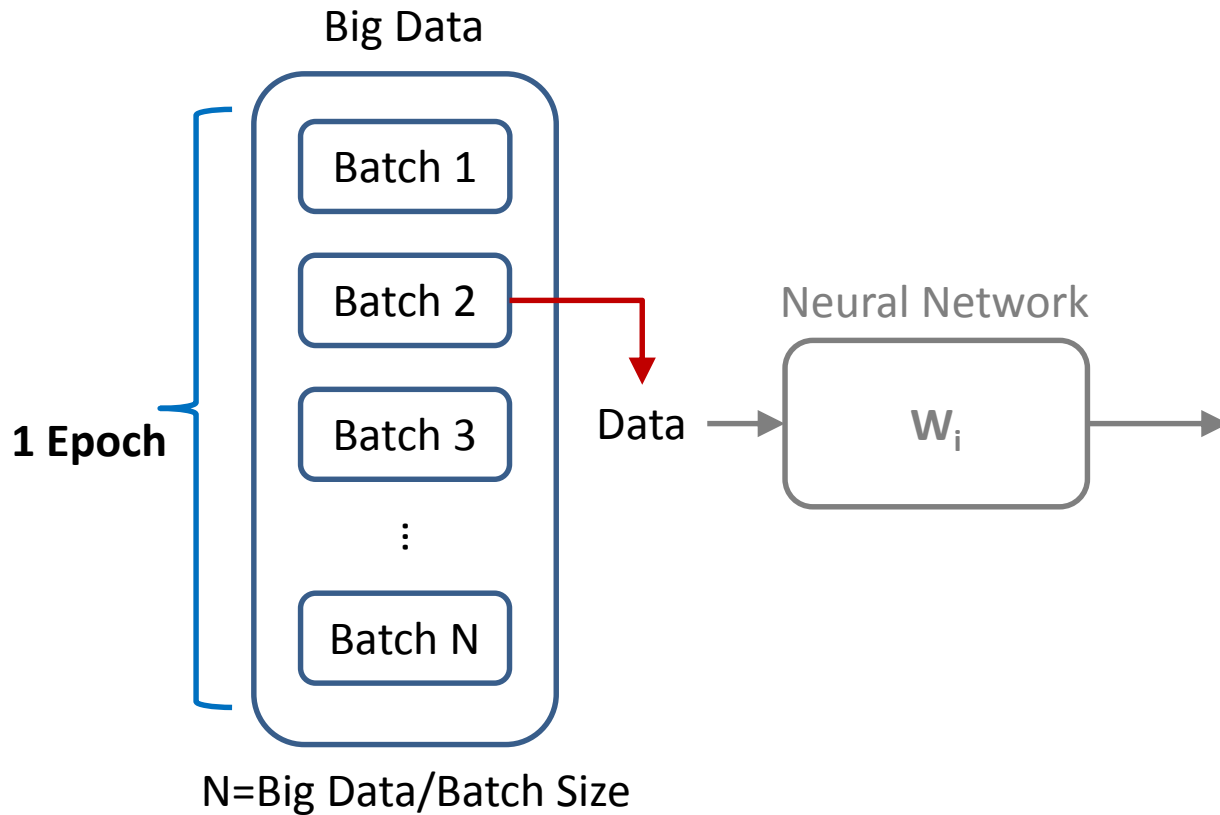
- Neural Network in Brief
- Concepts of PyTorch
- Multi-GPU Processing
- RNN
- Transfer Learning
- Comparison with TensorFlow

Neural Network in Brief

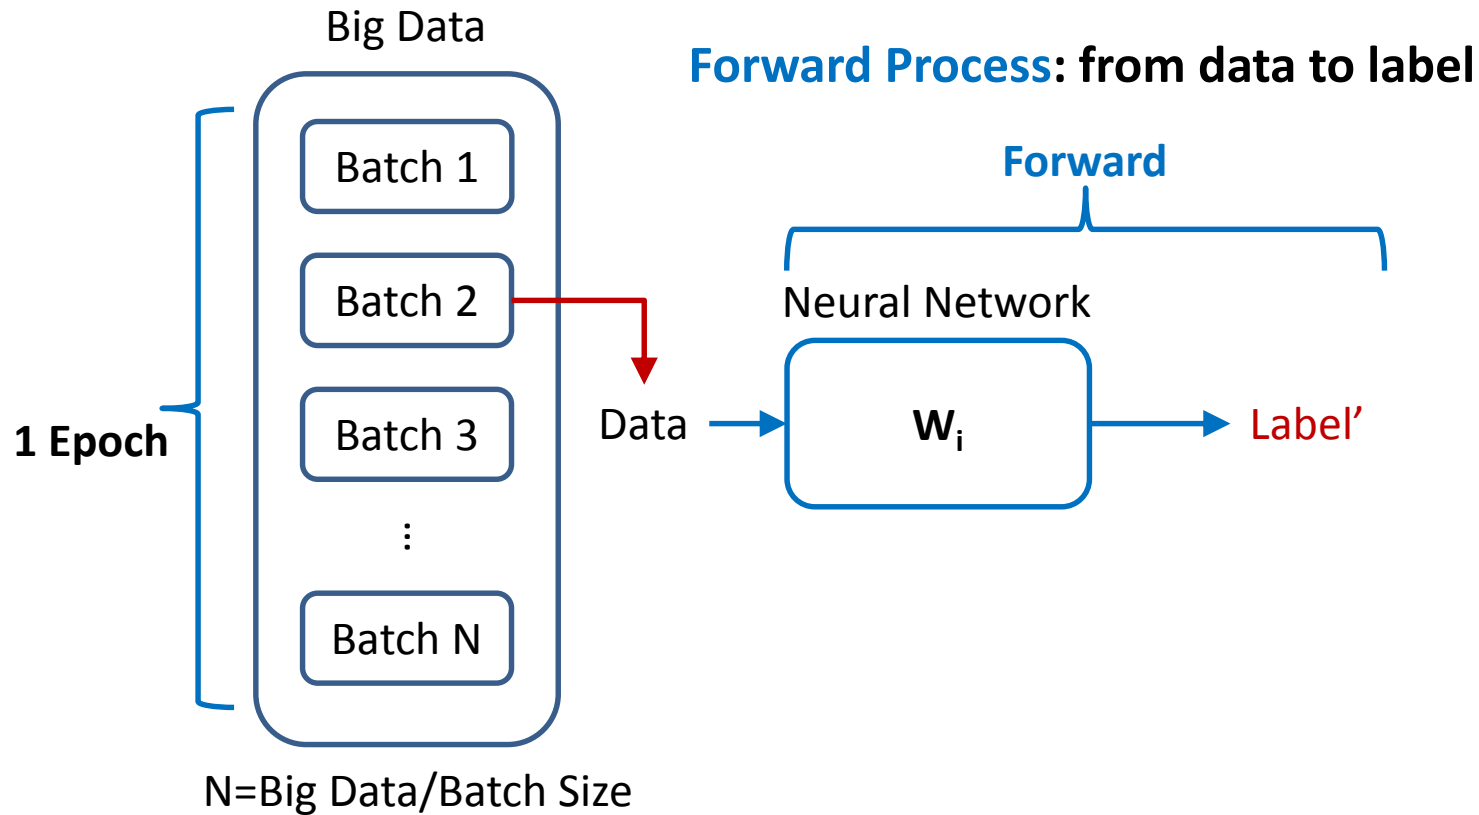
- Supervised Learning
 - Learning a function f , that $f(x)=y$

| Trying to learn $f(\cdot)$, that $f(x)=y$ | |
|--|-------|
| Data | Label |
| X1 | Y1 |
| X2 | Y2 |
| ... | ... |

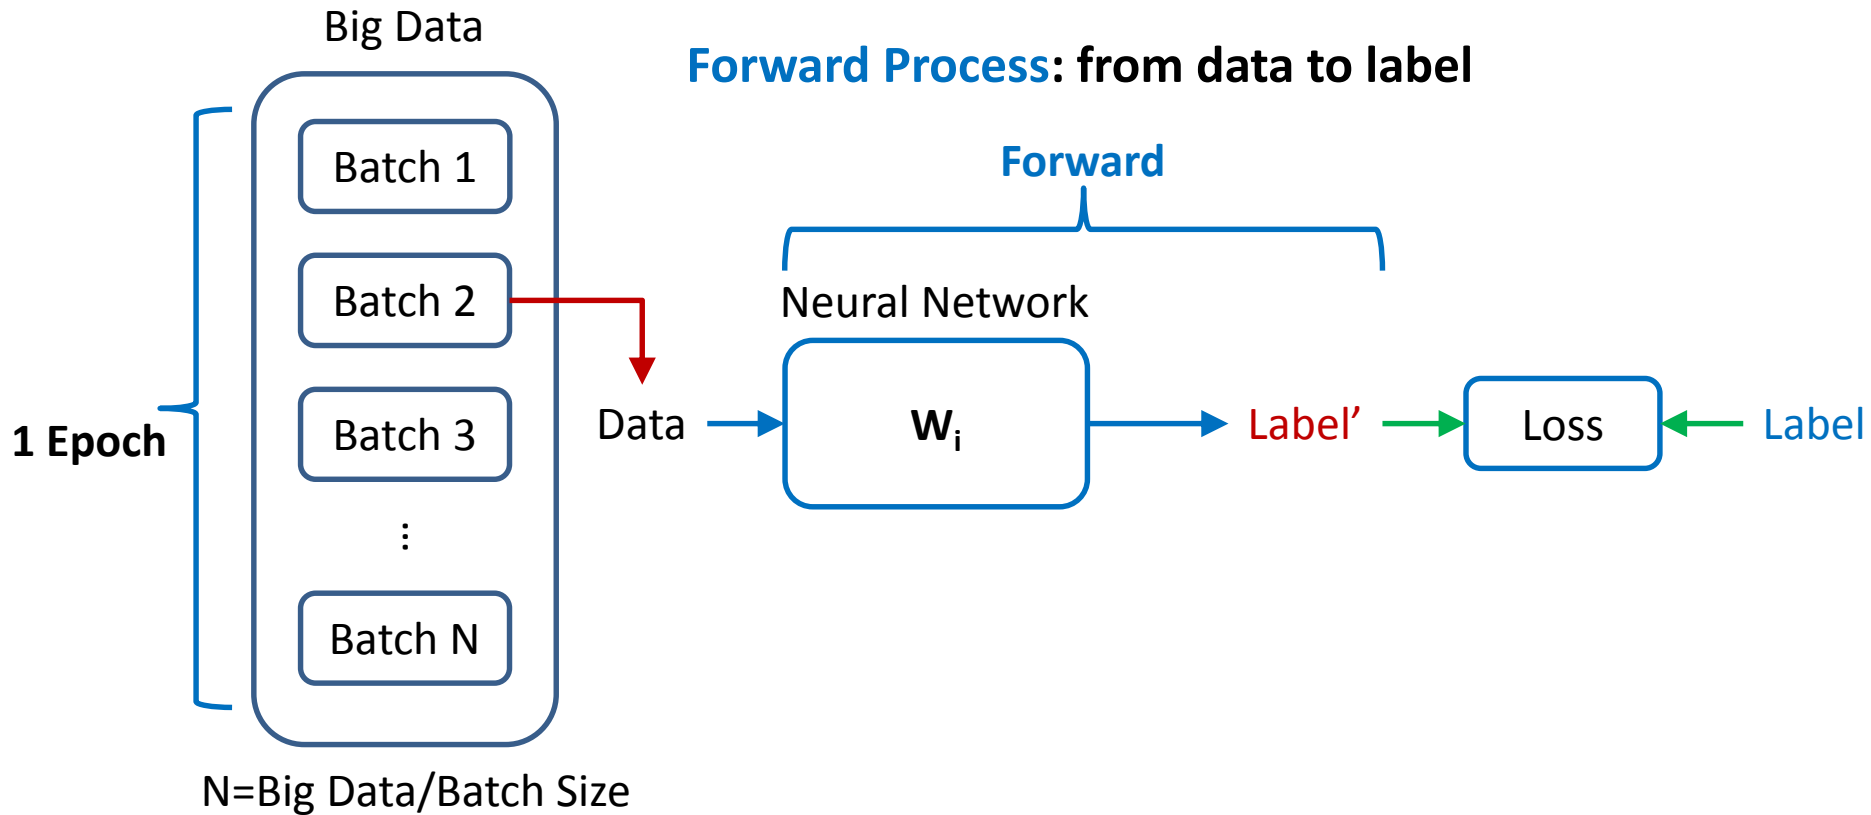
Neural Network in Brief



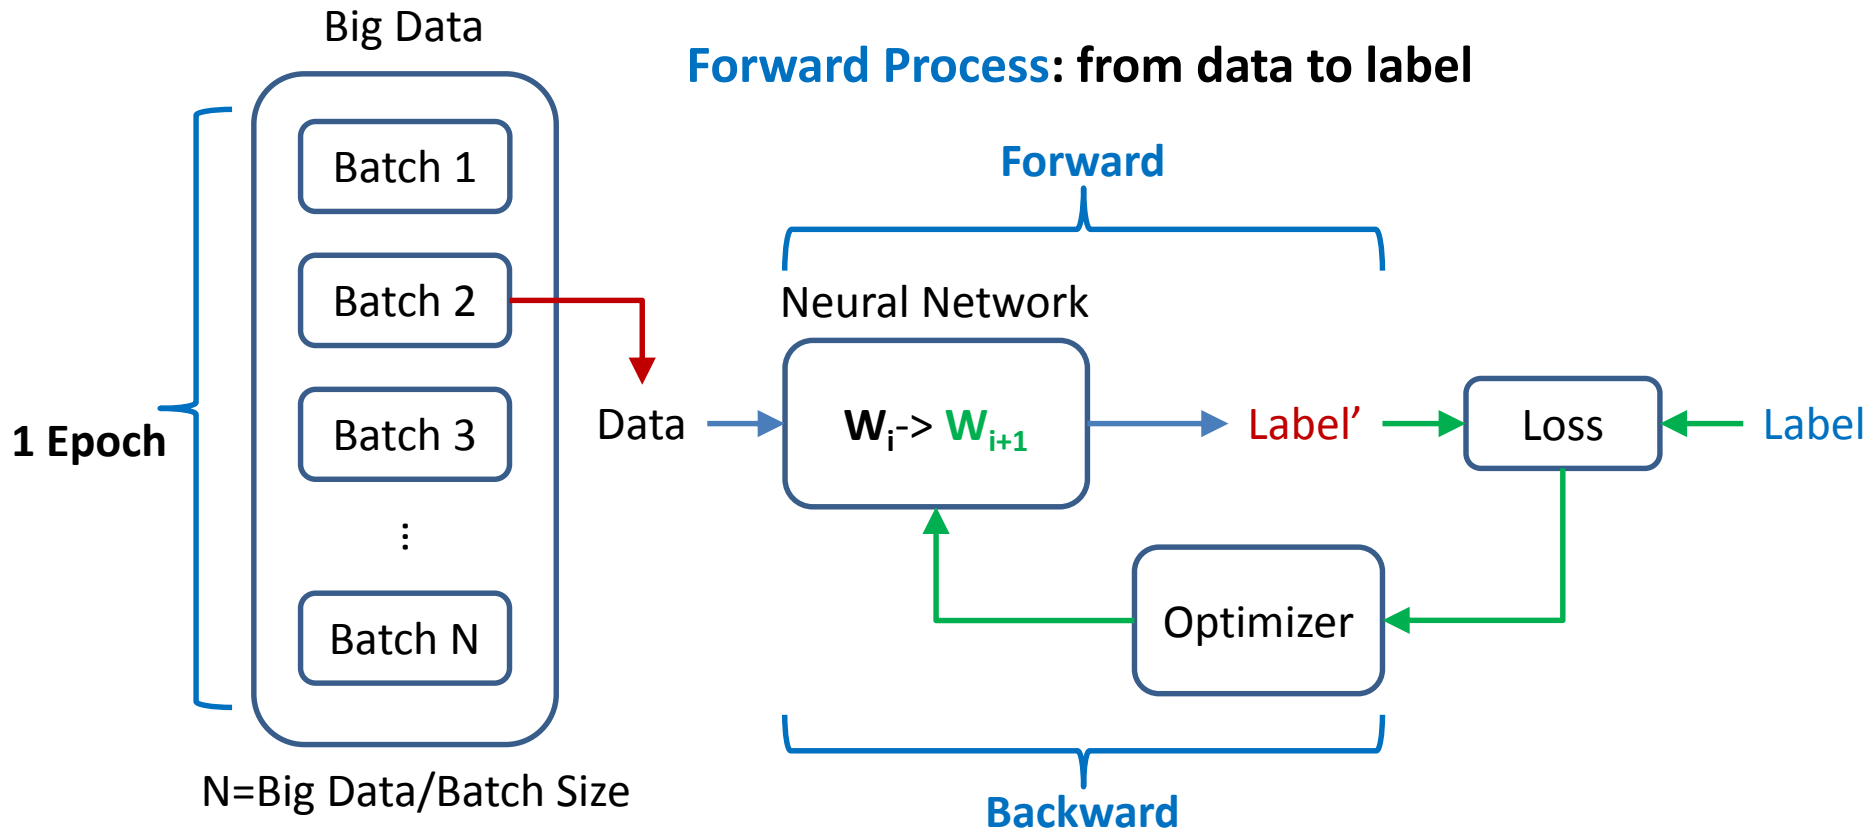
Neural Network in Brief



Neural Network in Brief

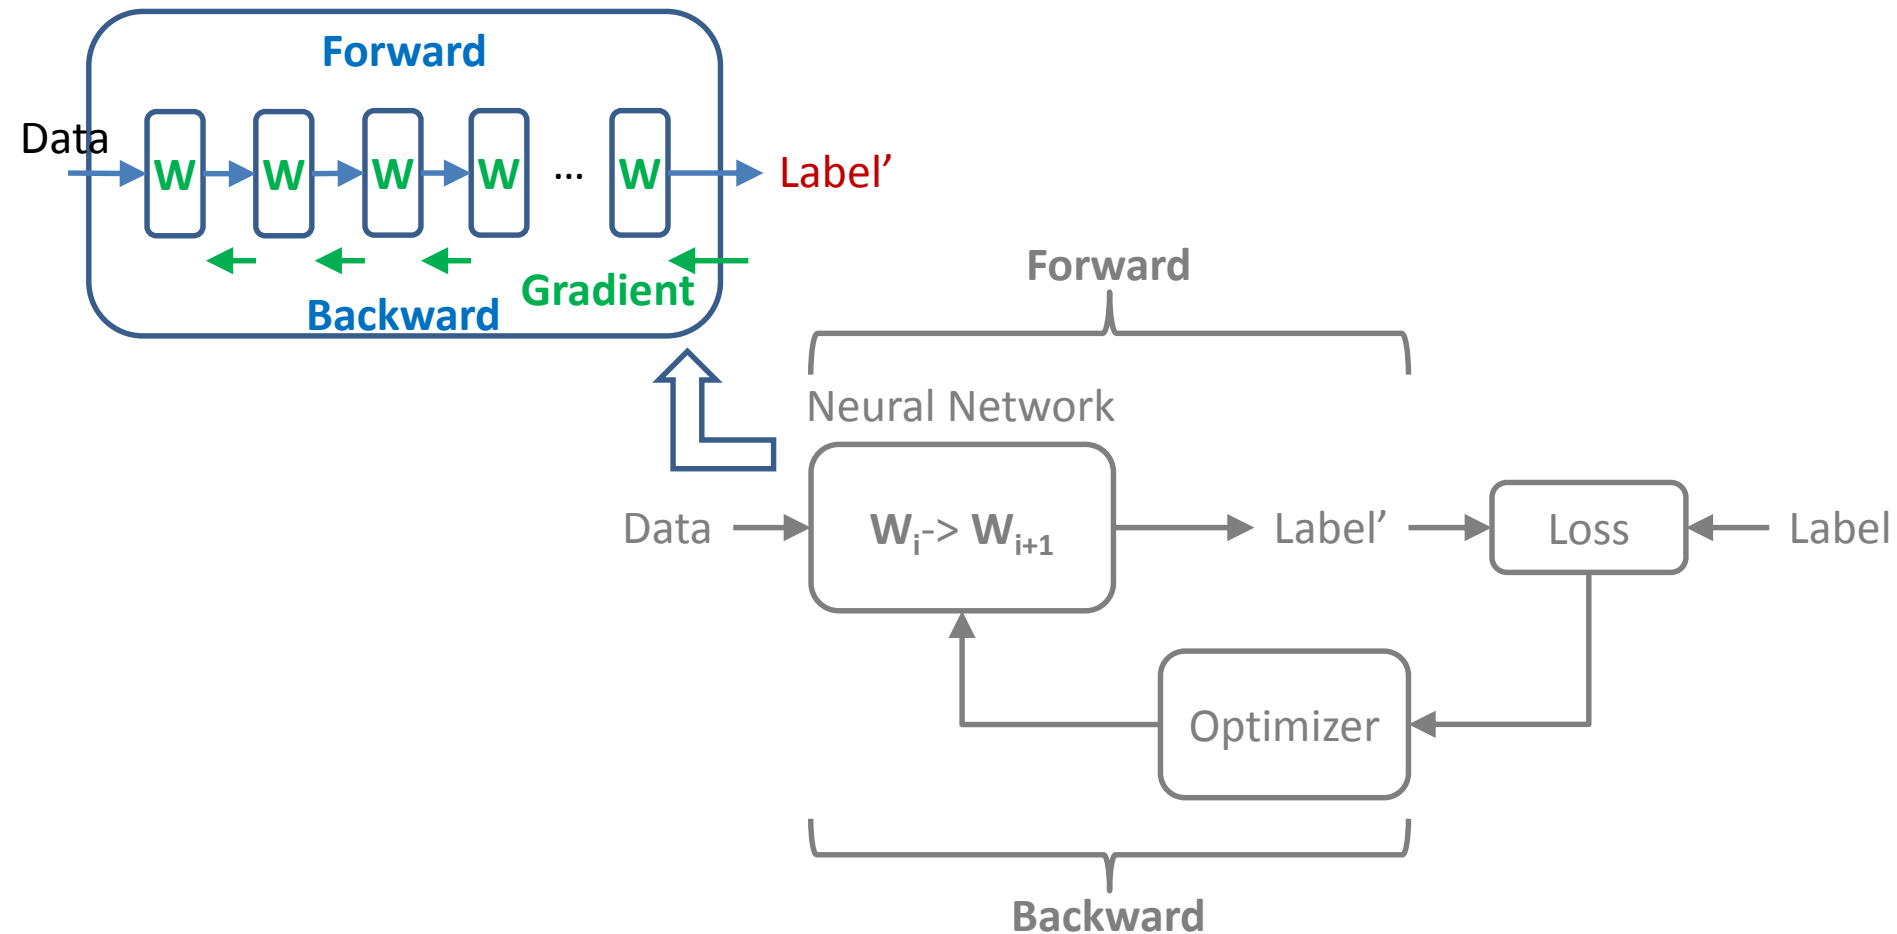


Neural Network in Brief



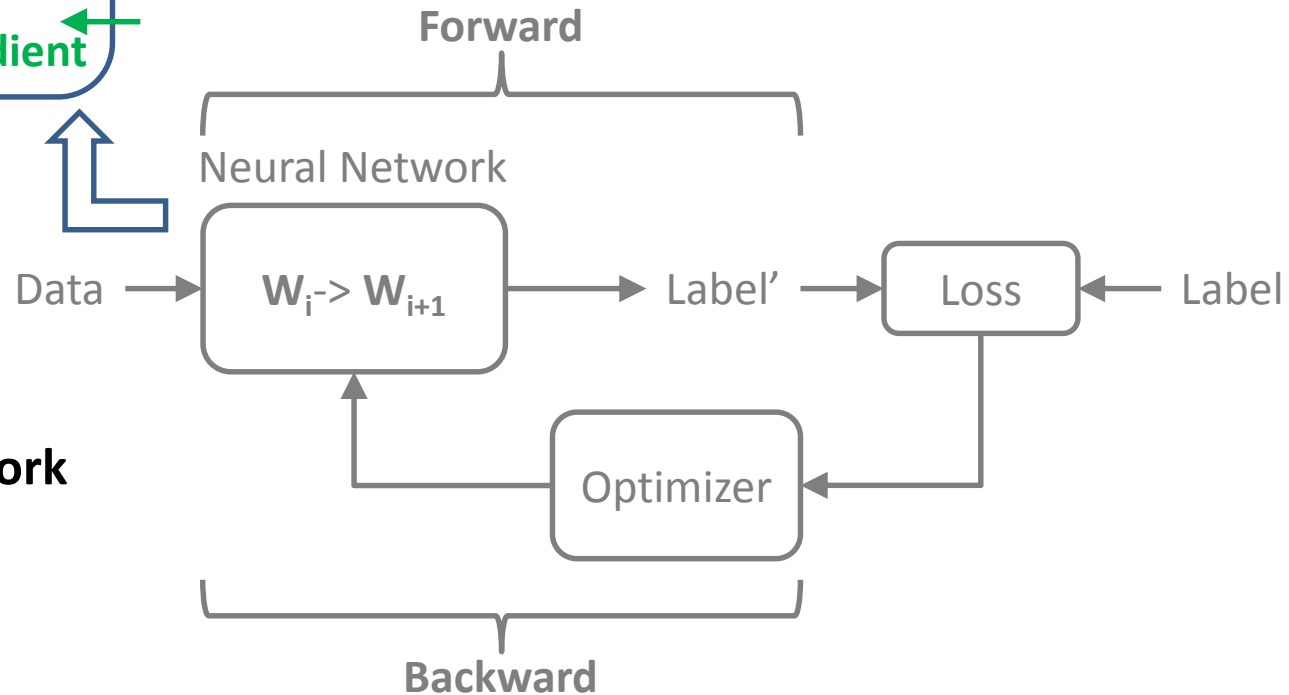

Neural Network in Brief

Inside the Neural Network



Neural Network in Brief

Inside the Neural Network



Data in the Neural Network

- Tensor (n-dim array)
- Gradient of Functions

Concepts of PyTorch

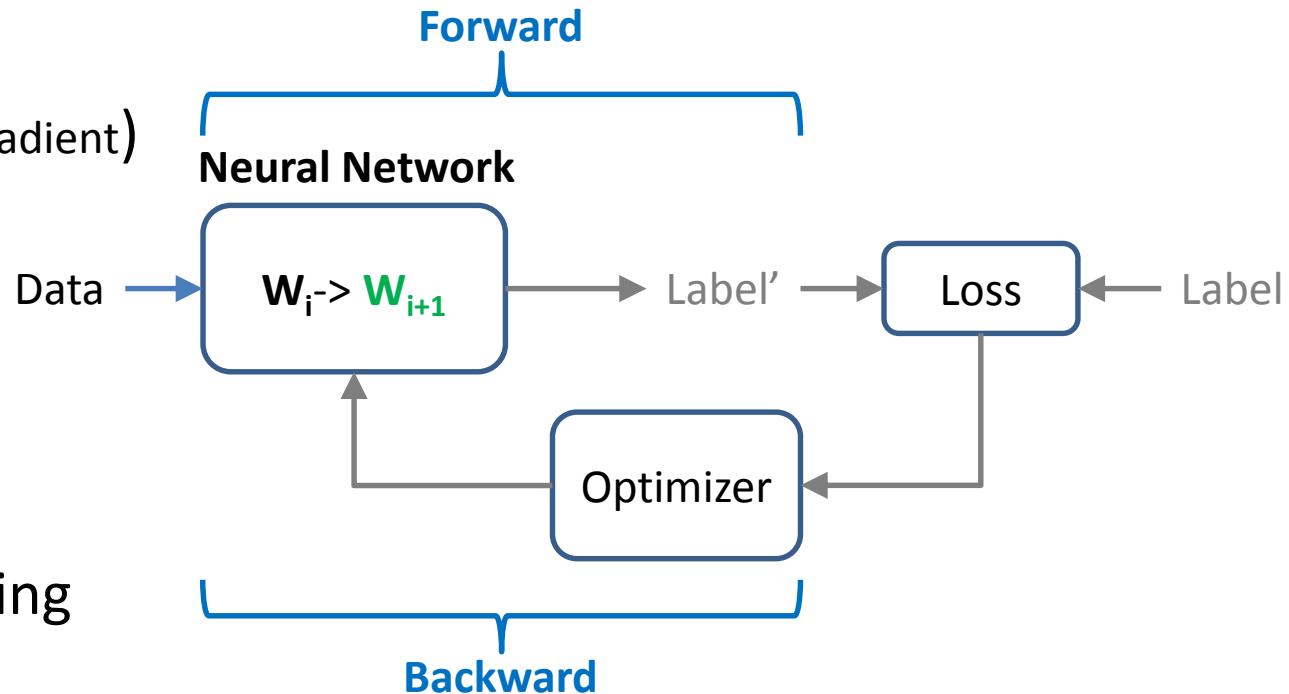
- Modules of PyTorch

Data:

- Tensor
- Variable (for Gradient)

Function:

- NN Modules
- Optimizer
- Loss Function
- Multi-Processing



Concepts of PyTorch

- Modules of PyTorch
- Similar to Numpy

Data:

- **Tensor**
- Variable (for Gradient)

Function:

- NN Modules
- Optimizer
- Loss Function
- Multi-Processing

```
x = torch.Tensor(5, 3)
print(x)
```

Out:

```
1.00000e-36 *
 0.0228  0.0000  1.3490
 0.0000  0.0958  0.0000
 0.0958  0.0000  0.0958
 0.0000  0.0958  0.0000
 0.0958  0.0000  0.0958
[torch.FloatTensor of size 5x3]
```

```
x = torch.rand(5, 3)
print(x)
```

Out:

```
0.2285  0.2843  0.1978
0.0092  0.8238  0.2703
0.1266  0.9613  0.2472
0.0918  0.2827  0.9803
0.9237  0.1946  0.0104
[torch.FloatTensor of size 5x3]
```

Concepts of PyTorch

- Modules of PyTorch

Data:

- **Tensor**
- Variable (for Gradient)

Function:

- NN Modules
- Optimizer
- Loss Function
- Multi-Processing

- Operations

- $z = x + y$
- `torch.add(x, y, out=z)`
- `y.add_(x)` # in-place

Concepts of PyTorch

- Modules of PyTorch
- Numpy Bridge

Data:

- **Tensor**
- Variable (for Gradient)

Function:

- NN Modules
- Optimizer
- Loss Function
- Multi-Processing

• To Numpy

- `a = torch.ones(5)`
- `b = a.numpy()`

• To Tensor

- `a = numpy.ones(5)`
- `b = torch.from_numpy(a)`

Concepts of PyTorch

- Modules of PyTorch
- CUDA Tensors

Data:

- **Tensor**
- Variable (for Gradient)

Function:

- NN Modules
- Optimizer
- Loss Function
- Multi-Processing

• Move to GPU

- `x = x.cuda()`
- `y = y.cuda()`
- `x+y`

Concepts of PyTorch

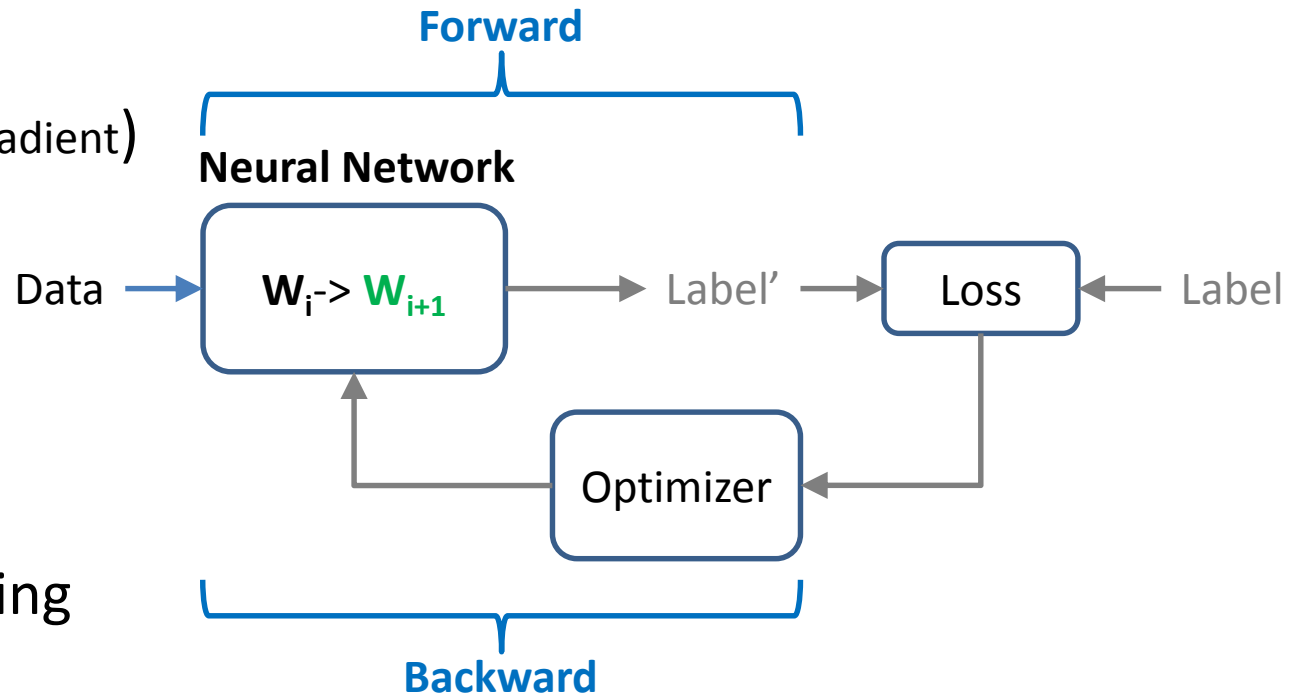
- Modules of PyTorch

Data:

- Tensor
- Variable (for Gradient)

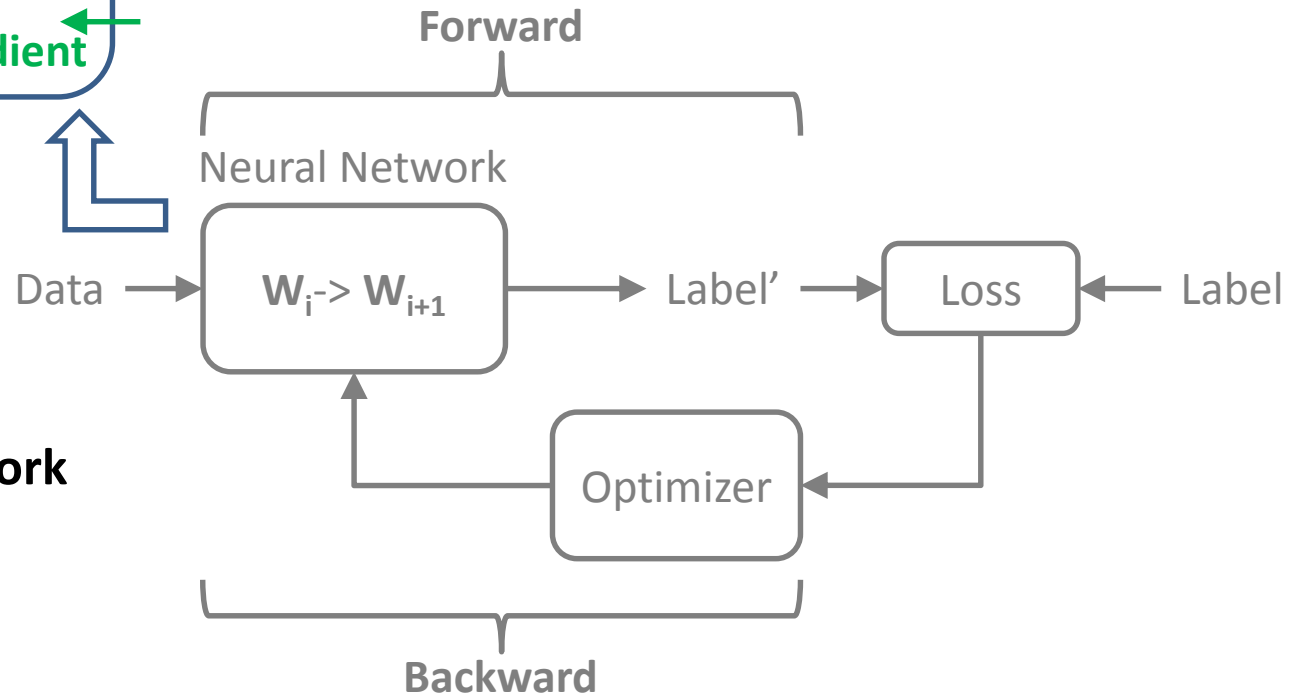

Function:

- NN Modules
- Optimizer
- Loss Function
- Multi-Processing



Neural Network in Brief

Inside the Neural Network



Data in the Neural Network

- Tensor (n-dim array)
- Gradient of Functions

Concepts of PyTorch

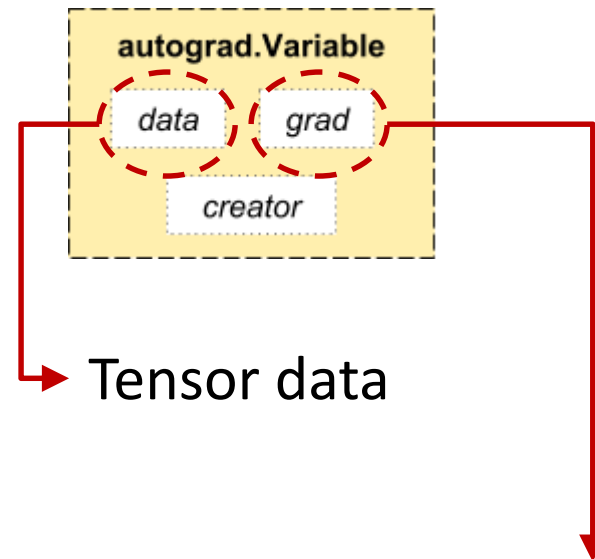
- Modules of PyTorch
- Variable

Data:

- Tensor
- **Variable** (for Gradient)

Function:

- NN Modules
- Optimizer
- Loss Function
- Multi-Processing



For **Current** Backward Process
Handled by PyTorch Automatically

Concepts of PyTorch

- Modules of PyTorch

Data:

- Tensor
- **Variable (for Gradient)**

Function:

- NN Modules
- Optimizer
- Loss Function
- Multi-Processing

- Variable

- `x = Variable(torch.ones(2, 2), requires_grad=True)`
- `print(x)`

Out: Variable containing:
1 1
1 1
[torch.FloatTensor of size 2x2]

- `y = x + 2`
- `z = y * y * 3`
- `out = z.mean()`
- `out.backward()`
- `print(x.grad)`

Out: Variable containing:
4.5000 4.5000
4.5000 4.5000
[torch.FloatTensor of size 2x2]

$$\text{out} = \frac{1}{4} \sum z_i$$

$$z_i = 3y_i^2 = 3(x_i + 2)^2$$

$$\frac{\partial \text{out}}{\partial x_i} = \frac{3}{2}(x_i + 2) = \frac{9}{2}$$



```

import torch
from torch.autograd import Variable
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):

    def __init__(self):
        super(Net, self).__init__()
        # 1 input image channel, 6 output channels, 5x5 square convolution
        # kernel
        self.conv1 = nn.Conv2d(1, 6, 5)
        self.conv2 = nn.Conv2d(6, 16, 5)
        # an affine operation: y = Wx + b
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        # Max pooling over a (2, 2) window
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
        # If the size is a square you can only specify a single number
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(-1, self.num_flat_features(x))
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

    def num_flat_features(self, x):
        size = x.size()[1:] # all dimensions except the batch dimension
        num_features = 1
        for s in size:
            num_features *= s
        return num_features

net = Net()
print(net)

```

- http://pytorch.org/tutorials/beginner/blitz/neural_networks_tutorial.html#define-the-network

Define modules
(must have)

```
import torch
from torch.autograd import Variable
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):

    def __init__(self):
        super(Net, self).__init__()
        # 1 input image channel, 6 output channels, 5x5 square convolution
        # kernel
        self.conv1 = nn.Conv2d(1, 6, 5)
        self.conv2 = nn.Conv2d(6, 16, 5)
        # an affine operation: y = Wx + b
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)
```

Build network
(must have)

```
    def forward(self, x):
        # Max pooling over a (2, 2) window
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
        # If the size is a square you can only specify a single number
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(-1, self.num_flat_features(x))
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

    def num_flat_features(self, x):
        size = x.size()[1:] # all dimensions except the batch dimension
        num_features = 1
        for s in size:
            num_features *= s
        return num_features
```

```
net = Net()
print(net)
```

- http://pytorch.org/tutorials/beginner/blitz/neural_networks_tutorial.html#define-the-network

Define modules
(must have)

```
import torch
from torch.autograd import Variable
import torch.nn as nn
import torch.nn.functional as F
```

```
class Net(nn.Module):
```

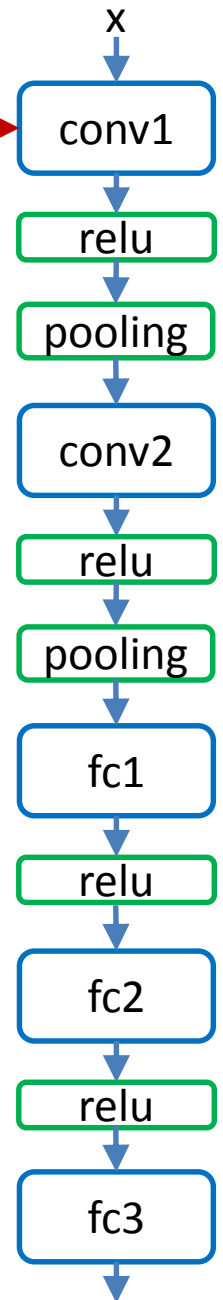
```
    def __init__(self):
        super(Net, self).__init__()
        # 1 input image channel, 6 output channels, 5x5 square convolution
        # kernel
        self.conv1 = nn.Conv2d(1, 6, 5)
        self.conv2 = nn.Conv2d(6, 16, 5)
        # an affine operation: y = Wx + b
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)
```

```
    def forward(self, x):
        # Max pooling over a (2, 2) window
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
        # If the size is a square you can only specify a single number
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(-1, self.num_flat_features(x))
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

```
    def num_flat_features(self, x):
        size = x.size()[1:] # all dimensions except the batch dimension
        num_features = 1
        for s in size:
            num_features *= s
        return num_features
```

```
net = Net()
print(net)
```

[Channel, H, W]: 1x32x32->6x28x28



Build network
(must have)

- http://pytorch.org/tutorials/beginner/blitz/neural_networks_tutorial.html#define-the-network

Define modules
(must have)

```
import torch
from torch.autograd import Variable
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):

    def __init__(self):
        super(Net, self).__init__()
        # 1 input image channel, 6 output channels, 5x5 square convolution
        # kernel
        self.conv1 = nn.Conv2d(1, 6, 5)
        self.conv2 = nn.Conv2d(6, 16, 5)
        # an affine operation: y = Wx + b
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)


    def forward(self, x):
        # Max pooling over a (2, 2) window
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
        # If the size is a square you can only specify a single number
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(-1, self.num_flat_features(x))
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

    def num_flat_features(self, x):
        size = x.size()[1:] # all dimensions except the batch dimension
        num_features = 1
        for s in size:
            num_features *= s
        return num_features

net = Net()
print(net)
```

Build network
(must have)

[Channel, H, W]: 6x28x28



- http://pytorch.org/tutorials/beginner/blitz/neural_networks_tutorial.html#define-the-network

Define modules
(must have)

```
import torch
from torch.autograd import Variable
import torch.nn as nn
import torch.nn.functional as F
```

```
class Net(nn.Module):
```

```
    def __init__(self):
        super(Net, self).__init__()
        # 1 input image channel, 6 output channels, 5x5 square convolution
        # kernel
        self.conv1 = nn.Conv2d(1, 6, 5)
        self.conv2 = nn.Conv2d(6, 16, 5)
        # an affine operation: y = Wx + b
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)
```

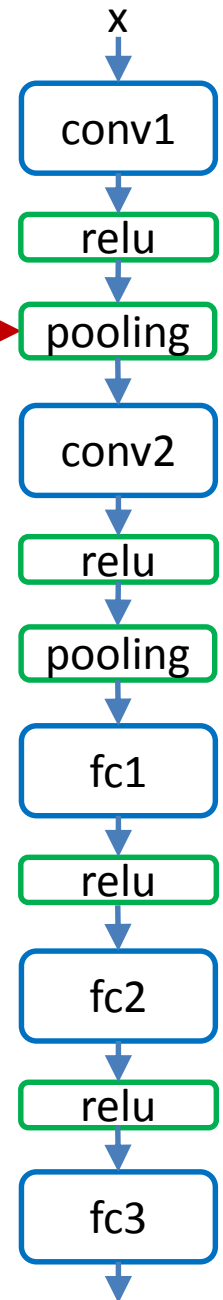
```
    def forward(self, x):
        # Max pooling over a (2, 2) window
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
        # If the size is a square you can only specify a single number
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(-1, self.num_flat_features(x))
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

```
    def num_flat_features(self, x):
        size = x.size()[1:] # all dimensions except the batch dimension
        num_features = 1
        for s in size:
            num_features *= s
        return num_features
```

```
net = Net()
print(net)
```

[Channel, H, W]: 6x28x28 -> 6x14x14

Build network
(must have)



- http://pytorch.org/tutorials/beginner/blitz/neural_networks_tutorial.html#define-the-network

Define modules
(must have)

```
import torch
from torch.autograd import Variable
import torch.nn as nn
import torch.nn.functional as F
```

```
class Net(nn.Module):
```

```
    def __init__(self):
```

```
        super(Net, self).__init__()
```

```
        # 1 input image channel, 6 output channels, 5x5 square convolution
```

```
        # kernel
```

```
        self.conv1 = nn.Conv2d(1, 6, 5) [Channel, H, W]: 6x14x14 -> 16x10x10
```

```
        self.conv2 = nn.Conv2d(6, 16, 5)
```

```
        # an affine operation:  $y = Wx + b$ 
```

```
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
```

```
        self.fc2 = nn.Linear(120, 84)
```

```
        self.fc3 = nn.Linear(84, 10)
```

```
    def forward(self, x):
```

```
        # Max pooling over a (2, 2) window
```

```
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
```

```
        # If the size is a square you can only specify a single number
```

```
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
```

```
        x = x.view(-1, self.num_flat_features(x))
```

```
        x = F.relu(self.fc1(x))
```

```
        x = F.relu(self.fc2(x))
```

```
        x = self.fc3(x)
```

```
        return x
```

```
    def num_flat_features(self, x):
```

```
        size = x.size()[1:] # all dimensions except the batch dimension
```

```
        num_features = 1
```

```
        for s in size:
```

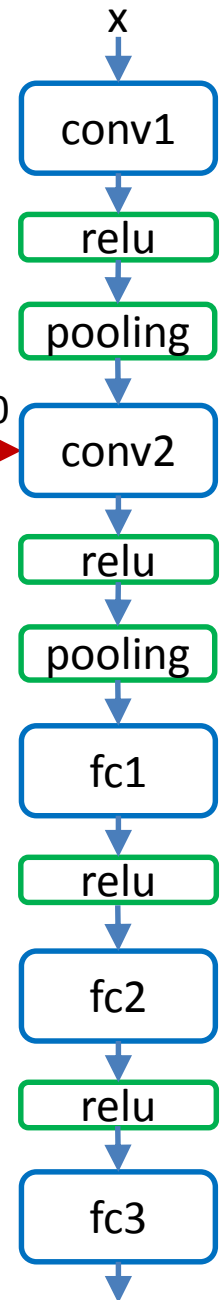
```
            num_features *= s
```

```
        return num_features
```

```
net = Net()
```

```
print(net)
```

Build network
(must have)



- http://pytorch.org/tutorials/beginner/blitz/neural_networks_tutorial.html#define-the-network

Define modules
(must have)

```
import torch
from torch.autograd import Variable
import torch.nn as nn
import torch.nn.functional as F
```

```
class Net(nn.Module):
```

```
    def __init__(self):
```

```
        super(Net, self).__init__()
```

```
        # 1 input image channel, 6 output channels, 5x5 square convolution
        # kernel
```

```
        self.conv1 = nn.Conv2d(1, 6, 5)
```

```
        self.conv2 = nn.Conv2d(6, 16, 5)
```

```
        # an affine operation:  $y = Wx + b$ 
```

```
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
```

```
        self.fc2 = nn.Linear(120, 84)
```

```
        self.fc3 = nn.Linear(84, 10)
```

[Channel, H, W]: 16x10x10

```
    def forward(self, x):
```

```
        # Max pooling over a (2, 2) window
```

```
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
```

```
        # If the size is a square you can only specify a single number
```

```
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
```

```
        x = x.view(-1, self.num_flat_features(x))
```

```
        x = F.relu(self.fc1(x))
```

```
        x = F.relu(self.fc2(x))
```

```
        x = self.fc3(x)
```

```
        return x
```

```
    def num_flat_features(self, x):
```

```
        size = x.size()[1:] # all dimensions except the batch dimension
```

```
        num_features = 1
```

```
        for s in size:
```


```
            num_features *= s
```

```
        return num_features
```

```
net = Net()
```

```
print(net)
```

Build network
(must have)



- http://pytorch.org/tutorials/beginner/blitz/neural_networks_tutorial.html#define-the-network

Define modules
(must have)

```
import torch
from torch.autograd import Variable
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):

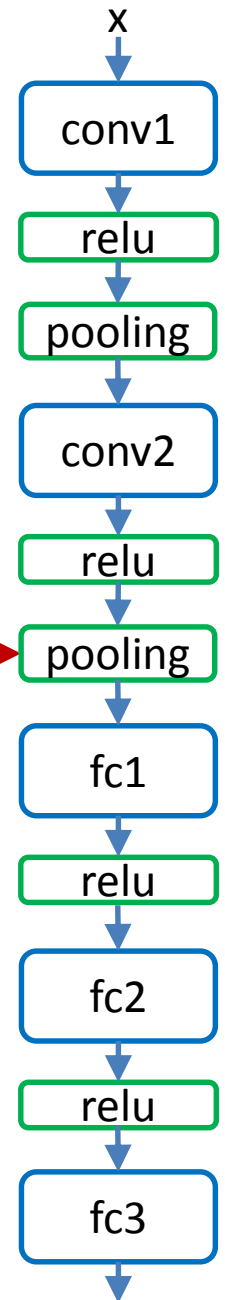
    def __init__(self):
        super(Net, self).__init__()
        # 1 input image channel, 6 output channels, 5x5 square convolution
        # kernel
        self.conv1 = nn.Conv2d(1, 6, 5)
        self.conv2 = nn.Conv2d(6, 16, 5)
        # an affine operation: y = Wx + b
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)
```

Build network
(must have)

```
    def forward(self, x):
        # Max pooling over a (2, 2) window
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
        # If the size is a square you can only specify a single number
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(-1, self.num_flat_features(x))
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

    def num_flat_features(self, x):
        size = x.size()[1:] # all dimensions except the batch dimension
        num_features = 1
        for s in size:
            num_features *= s
        return num_features
```

```
net = Net()
print(net)
```



- http://pytorch.org/tutorials/beginner/blitz/neural_networks_tutorial.html#define-the-network

Define modules
(must have)

```
import torch
from torch.autograd import Variable
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):

    def __init__(self):
        super(Net, self).__init__()
        # 1 input image channel, 6 output channels, 5x5 square convolution
        # kernel
        self.conv1 = nn.Conv2d(1, 6, 5)
        self.conv2 = nn.Conv2d(6, 16, 5)
        # an affine operation: y = Wx + b
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)
```

Build network
(must have)

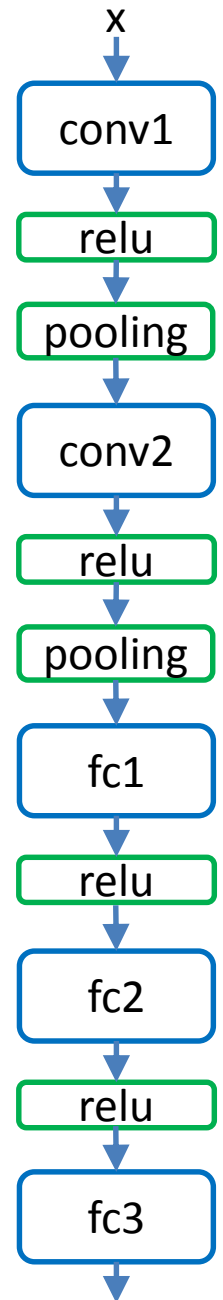
```
    def forward(self, x):
        # Max pooling over a (2, 2) window
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
        # If the size is a square you can only specify a single number
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(-1, self.num_flat_features(x))
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

```
    def num_flat_features(self, x):
        size = x.size()[1:] # all dimensions except the batch dimension
        num_features = 1
        for s in size:
            num_features *= s
        return num_features
```

```
net = Net()
print(net)
```

Tensor: [Batch N, Channel, H, W]

Flatten the Tensor
16x5x5



- http://pytorch.org/tutorials/beginner/blitz/neural_networks_tutorial.html#define-the-network

Define modules
(must have)

```
import torch
from torch.autograd import Variable
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):

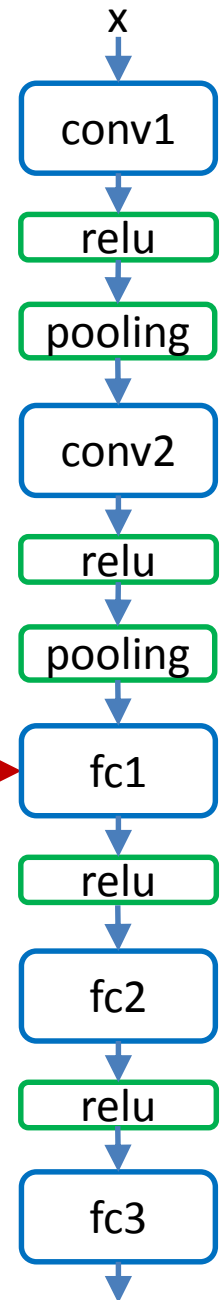
    def __init__(self):
        super(Net, self).__init__()
        # 1 input image channel, 6 output channels, 5x5 square convolution
        # kernel
        self.conv1 = nn.Conv2d(1, 6, 5)
        self.conv2 = nn.Conv2d(6, 16, 5)
        # an affine operation: y = Wx + b
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        # Max pooling over a (2, 2) window
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
        # If the size is a square you can only specify a single number
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(-1, self.num_flat_features(x))
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

    def num_flat_features(self, x):
        size = x.size()[1:] # all dimensions except the batch dimension
        num_features = 1
        for s in size:
            num_features *= s
        return num_features

net = Net()
print(net)
```

Build network
(must have)



- http://pytorch.org/tutorials/beginner/blitz/neural_networks_tutorial.html#define-the-network

Define modules
(must have)

```
import torch
from torch.autograd import Variable
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):

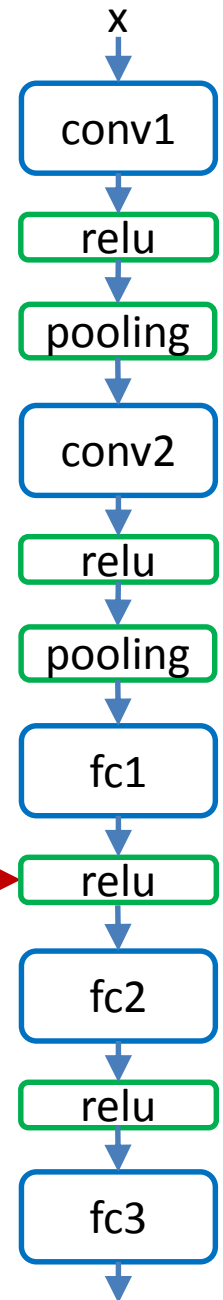
    def __init__(self):
        super(Net, self).__init__()
        # 1 input image channel, 6 output channels, 5x5 square convolution
        # kernel
        self.conv1 = nn.Conv2d(1, 6, 5)
        self.conv2 = nn.Conv2d(6, 16, 5)
        # an affine operation: y = Wx + b
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        # Max pooling over a (2, 2) window
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
        # If the size is a square you can only specify a single number
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(-1, self.num_flat_features(x))
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

    def num_flat_features(self, x):
        size = x.size()[1:] # all dimensions except the batch dimension
        num_features = 1
        for s in size:
            num_features *= s
        return num_features

net = Net()
print(net)
```

Build network
(must have)



- http://pytorch.org/tutorials/beginner/blitz/neural_networks_tutorial.html#define-the-network

Define modules
(must have)

```
import torch
from torch.autograd import Variable
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):


    def __init__(self):
        super(Net, self).__init__()
        # 1 input image channel, 6 output channels, 5x5 square convolution
        # kernel
        self.conv1 = nn.Conv2d(1, 6, 5)
        self.conv2 = nn.Conv2d(6, 16, 5)
        # an affine operation: y = Wx + b
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        # Max pooling over a (2, 2) window
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
        # If the size is a square you can only specify a single number
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(-1, self.num_flat_features(x))
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

    def num_flat_features(self, x):
        size = x.size()[1:] # all dimensions except the batch dimension
        num_features = 1
        for s in size:
            num_features *= s
        return num_features

net = Net()
print(net)
```

Build network
(must have)



- http://pytorch.org/tutorials/beginner/blitz/neural_networks_tutorial.html#define-the-network

Define modules
(must have)

```
import torch
from torch.autograd import Variable
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):


    def __init__(self):
        super(Net, self).__init__()
        # 1 input image channel, 6 output channels, 5x5 square convolution
        # kernel
        self.conv1 = nn.Conv2d(1, 6, 5)
        self.conv2 = nn.Conv2d(6, 16, 5)
        # an affine operation: y = Wx + b
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        # Max pooling over a (2, 2) window
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
        # If the size is a square you can only specify a single number
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(-1, self.num_flat_features(x))
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

    def num_flat_features(self, x):
        size = x.size()[1:] # all dimensions except the batch dimension
        num_features = 1
        for s in size:
            num_features *= s
        return num_features

net = Net()
print(net)
```

Build network
(must have)



- http://pytorch.org/tutorials/beginner/blitz/neural_networks_tutorial.html#define-the-network

Define modules
(must have)

```
import torch
from torch.autograd import Variable
import torch.nn as nn
import torch.nn.functional as F

class Net(nn.Module):

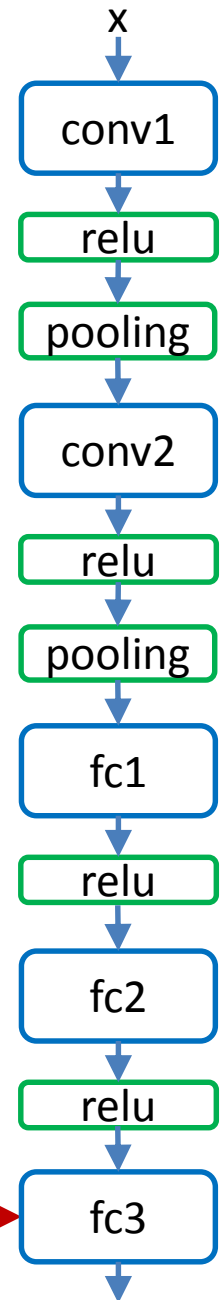
    def __init__(self):
        super(Net, self).__init__()
        # 1 input image channel, 6 output channels, 5x5 square convolution
        # kernel
        self.conv1 = nn.Conv2d(1, 6, 5)
        self.conv2 = nn.Conv2d(6, 16, 5)
        # an affine operation: y = Wx + b
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        # Max pooling over a (2, 2) window
        x = F.max_pool2d(F.relu(self.conv1(x)), (2, 2))
        # If the size is a square you can only specify a single number
        x = F.max_pool2d(F.relu(self.conv2(x)), 2)
        x = x.view(-1, self.num_flat_features(x))
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x

    def num_flat_features(self, x):
        size = x.size()[1:] # all dimensions except the batch dimension
        num_features = 1
        for s in size:
            num_features *= s
        return num_features

net = Net()
print(net)
```

Build network
(must have)



- http://pytorch.org/tutorials/beginner/blitz/neural_networks_tutorial.html#define-the-network

Concepts of PyTorch

- Modules of PyTorch

Data:

- Tensor
- Variable (for Gradient)

Function:

- **NN Modules**
- Optimizer
- Loss Function
- Multi-Processing

- NN Modules (torch.nn)

- Modules built on Variable
- Gradient handled by PyTorch

- Common Modules

- Convolution layers
- Linear layers
- Pooling layers
- Dropout layers
- Etc...

NN Modules

- Convolution Layer

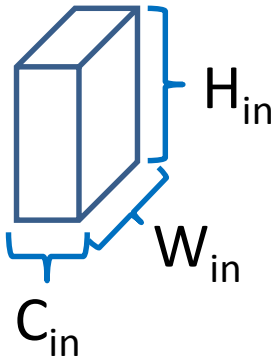
- N-th Batch (N), Channel (C)
- torch.nn.Conv1d: input [N, C, W] # moving kernel in 1D
- torch.nn.Conv2d: input [N, C, H, W] # moving kernel in 2D
- torch.nn.Conv3d: input [N, C, D, H, W] # moving kernel in 3D
- Example:
- torch.nn.conv2d(in_channels=3, out_channels=16, kernel_size=3, padding=1)

NN Modules

- Convolution Layer

- N-th Batch (N), Channel (C)
- `torch.nn.Conv1d`: input [N, C, W] # moving kernel in 1D
- `torch.nn.Conv2d`: input [N, C, H, W] # moving kernel in 2D
- `torch.nn.Conv3d`: input [N, C, D, H, W] # moving kernel in 3D

Input for Conv2d



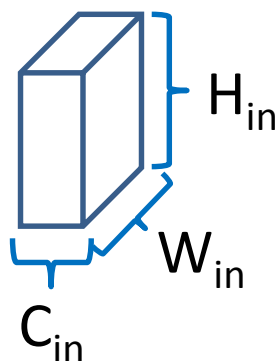
*: convolution

NN Modules


- Convolution Layer

- N-th Batch (N), Channel (C)
- `torch.nn.Conv1d`: input [N, C, W] # moving kernel in 1D
- `torch.nn.Conv2d`: input [N, C, H, W] # moving kernel in 2D
- `torch.nn.Conv3d`: input [N, C, D, H, W] # moving kernel in 3D

Input for Conv2d



1st kernel



*

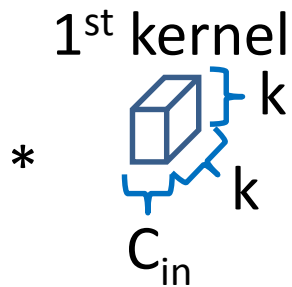

∗: convolution

NN Modules

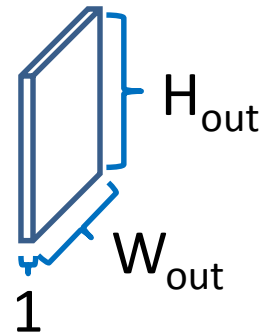
- Convolution Layer

- N-th Batch (N), Channel (C)
- `torch.nn.Conv1d`: input [N, C, W] # moving kernel in 1D
- `torch.nn.Conv2d`: input [N, C, H, W] # moving kernel in 2D
- `torch.nn.Conv3d`: input [N, C, D, H, W] # moving kernel in 3D

Input for Conv2d



=



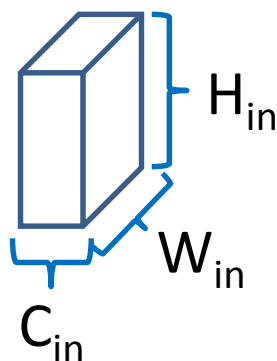
*: convolution

NN Modules

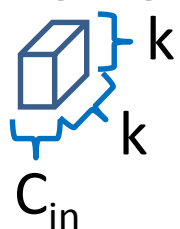
- Convolution Layer

- N-th Batch (N), Channel (C)
- torch.nn.Conv1d: input [N, C, W] # moving kernel in 1D
- torch.nn.Conv2d: input [N, C, H, W] # moving kernel in 2D
- torch.nn.Conv3d: input [N, C, D, H, W] # moving kernel in 3D


Input for Conv2d



1st kernel



=



$$W_{\text{out}} = \text{floor}\left(\frac{W_{\text{in}} + 2 \times \text{padding} - \text{dilation} \times (k - 1) - 1}{\text{stride}} + 1\right)$$

p=1 d=1 k=3
s=1, moving step size

$$W_{\text{out}} = \text{floor}\left(\frac{W_{\text{in}} + 2 \times 1 - 1 \times (3 - 1) - 1}{1} + 1\right) = W_{\text{in}}$$


∗: convolution

NN Modules


- Convolution Layer

- N-th Batch (N), Channel (C)
- torch.nn.Conv1d: input [N, C, W] # moving kernel in 1D
- torch.nn.Conv2d: input [N, C, H, W] # moving kernel in 2D
- torch.nn.Conv3d: input [N, C, D, H, W] # moving kernel in 3D


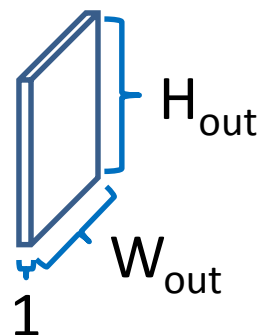
Input for Conv2d



1st kernel



=



$$W_{\text{out}} = \text{floor}\left(\frac{W_{\text{in}} + 2 \times \text{padding} - \text{dilation} \times (k - 1) - 1}{\text{stride}} + 1\right)$$

p=1, d=1, k=3, s=1, moving step size


$$W_{\text{out}} = \text{floor}\left(\frac{W_{\text{in}} + 2 \times 1 - 1 \times (3 - 1) - 1}{1} + 1\right) = W_{\text{in}}$$

∗: convolution

NN Modules

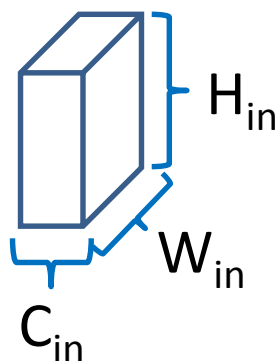
- Convolution Layer

- N-th Batch (N), Channel (C)
- torch.nn.Conv1d: input [N, C, W] # moving kernel in 1D
- torch.nn.Conv2d: input [N, C, H, W] # moving kernel in 2D
- torch.nn.Conv3d: input [N, C, D, H, W] # moving kernel in 3D




Input for Conv2d


1st kernel



*



=



$$W_{\text{out}} = \text{floor}\left(\frac{W_{\text{in}} + 2 \times \text{padding} - \text{dilation} \times (k - 1) - 1}{\text{stride}} + 1\right)$$

p=1, d=1, k=3, s=1, moving step size

$$W_{\text{out}} = \text{floor}\left(\frac{W_{\text{in}} + 2 \times 1 - 1 \times (3 - 1) - 1}{1} + 1\right) = W_{\text{in}}$$


∗: convolution

NN Modules


• Convolution Layer

- N-th Batch (N), Channel (C)
- torch.nn.Conv1d: input [N, C, W] # moving kernel in 1D
- torch.nn.Conv2d: input [N, C, H, W] # moving kernel in 2D
- torch.nn.Conv3d: input [N, C, D, H, W] # moving kernel in 3D



Input for Conv2d



1st kernel



=



$$W_{\text{out}} = \text{floor}\left(\frac{W_{\text{in}} + 2 \times \text{padding} - \text{dilation} \times (k - 1) - 1}{\text{stride}} + 1\right)$$

s=1, moving step size

$$W_{\text{out}} = \text{floor}\left(\frac{W_{\text{in}} + 2 \times 1 - 1 \times (3 - 1) - 1}{1} + 1\right) = W_{\text{in}}$$


*: convolution

NN Modules

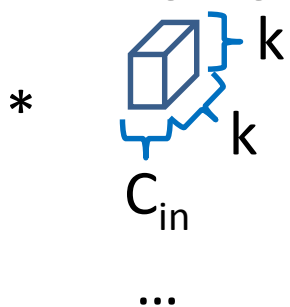
- Convolution Layer

- N-th Batch (N), Channel (C)
- `torch.nn.Conv1d`: input [N, C, W] # moving kernel in 1D
- `torch.nn.Conv2d`: input [N, C, H, W] # moving kernel in 2D
- `torch.nn.Conv3d`: input [N, C, D, H, W] # moving kernel in 3D


Input for Conv2d




1st kernel




=



C_{out} -th kernel



=



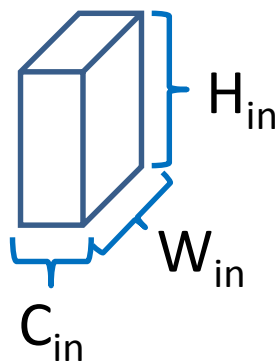
*: convolution

NN Modules


- Convolution Layer

- N-th Batch (N), Channel (C)
- `torch.nn.Conv1d`: input [N, C, W] # moving kernel in 1D
- `torch.nn.Conv2d`: input [N, C, H, W] # moving kernel in 2D
- `torch.nn.Conv3d`: input [N, C, D, H, W] # moving kernel in 3D


Input for Conv2d



1st kernel




=





...

...

C_{out} -th kernel



=



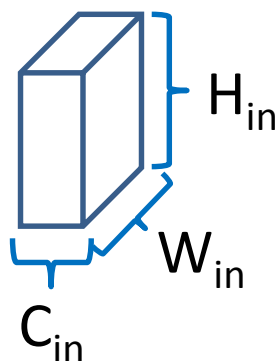
*: convolution

NN Modules

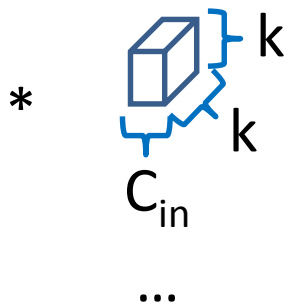
- Convolution Layer

- N-th Batch (N), Channel (C)
- `torch.nn.Conv1d`: input [N, C, W] # moving kernel in 1D
- `torch.nn.Conv2d`: input [N, C, H, W] # moving kernel in 2D
- `torch.nn.Conv3d`: input [N, C, D, H, W] # moving kernel in 3D

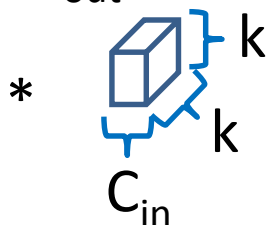
Input for Conv2d



1st kernel



C_{out} -th kernel




of parameters

$$O(C_{in} \times k^2 \times C_{out})$$

*: convolution


NN Modules

- Linear Layer
 - `torch.nn.Linear(in_features=3, out_features=5)`
 - $y = Ax + b$



NN Modules


- Dropout Layer
 - `torch.nn.Dropout(p)`
 - Random zeros the input with probability p
 - Output are scaled by $1/(1-p)$



NN Modules

- Pooling Layer

- `torch.nn.AvgPool2d(kernel_size=2, stride=2, padding=0)`
- `torch.nn.MaxPool2d(kernel_size=2, stride=2, padding=0)`



$$W_{\text{out}} = \text{floor}\left(\frac{W_{\text{in}} + 2 \times \overset{\text{p}=0}{\text{padding}} - \overset{\text{d}=1}{\text{dilation}} \times (\overset{\text{k}=2}{k} - 1) - 1}{\underset{\text{s}=2, \text{ moving step size}}{\text{stride}}} + 1\right)$$

$$W_{\text{out}} = \text{floor}\left(\frac{W_{\text{in}} + 2 \times 0 - 1 \times (2 - 1) - 1}{2} + 1\right) = \frac{W_{\text{in}}}{2}$$

Concepts of PyTorch

- Modules of PyTorch

Data:

- Tensor
- Variable (for Gradient)

Function:

- **NN Modules**
- Optimizer
- Loss Function
- Multi-Processing

- NN Modules (torch.nn)

- Modules built on Variable
- Gradient handled by PyTorch

- Common Modules

- Convolution layers
- Linear layers
- Pooling layers
- Dropout layers
- Etc...

Concepts of PyTorch

- Modules of PyTorch

Data:

- Tensor
- Variable (for Gradient)

Function:

- NN Modules
- **Optimizer**
- **Loss Function**
- Multi-Processing

- Optimizer (torch.optim)

- SGD
- Adagrad
- Adam
- RMSprop
- ...
- 9 Optimizers (PyTorch 0.2)

- Loss (torch.nn)

- L1Loss
- MSELoss
- CrossEntropy
- ...
- 18 Loss Functions (PyTorch 0.2)

What We Build?

Define modules
(must have)

```
# -*- coding: utf-8 -*-
import torch
from torch.autograd import Variable

class TwoLayerNet(torch.nn.Module):
    def __init__(self, D_in, H, D_out):
        """
        In the constructor we instantiate two nn.Linear modules and assign them as
        member variables.
        """
        super(TwoLayerNet, self).__init__()
        self.linear1 = torch.nn.Linear(D_in, H)
        self.linear2 = torch.nn.Linear(H, D_out)

    def forward(self, x):
        """
        In the forward function we accept a Variable of input data and we must return
        a Variable of output data. We can use Modules defined in the constructor as
        well as arbitrary operators on Variables.
        """
        h_relu = self.linear1(x).clamp(min=0)
        y_pred = self.linear2(h_relu)
        return y_pred

# N is batch size; D_in is input dimension;
# H is hidden dimension; D_out is output dimension.
N, D_in, H, D_out = 64, 1000, 100, 10

# Create random Tensors to hold inputs and outputs, and wrap them in Variables
x = Variable(torch.randn(N, D_in))
y = Variable(torch.randn(N, D_out), requires_grad=False)

# Construct our model by instantiating the class defined above
model = TwoLayerNet(D_in, H, D_out)


# Construct our loss function and an Optimizer. The call to model.parameters()
# in the SGD constructor will contain the learnable parameters of the two
# nn.Linear modules which are members of the model.
criterion = torch.nn.MSELoss(size_average=False)
optimizer = torch.optim.SGD(model.parameters(), lr=1e-4)
for t in range(500):
    # Forward pass: Compute predicted y by passing x to the model
    y_pred = model(x)

    # Compute and print loss
    loss = criterion(y_pred, y)
    print(t, loss.data[0])

    # Zero gradients, perform a backward pass, and update the weights.
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
```

Build network
(must have)

What We Build?



```
# -*- coding: utf-8 -*-
import torch
from torch.autograd import Variable
```

```
class TwoLayerNet(torch.nn.Module):
    def __init__(self, D_in, H, D_out):
        """
        In the constructor we instantiate two nn.Linear modules and assign them as
        member variables.
        """
        super(TwoLayerNet, self).__init__()
        self.linear1 = torch.nn.Linear(D_in, H)
        self.linear2 = torch.nn.Linear(H, D_out)
```

```
    def forward(self, x):
        """
        In the forward function we accept a Variable of input data and we must return
        a Variable of output data. We can use Modules defined in the constructor as
        well as arbitrary operators on Variables.
        """
        h_relu = self.linear1(x).clamp(min=0)
        y_pred = self.linear2(h_relu)
        return y_pred
```

```
# N is batch size; D_in is input dimension;
# H is hidden dimension; D_out is output dimension.
N, D_in, H, D_out = 64, 1000, 100, 10
```

```
# Create random Tensors to hold inputs and outputs, and wrap them in Variables
x = Variable(torch.randn(N, D_in))
y = Variable(torch.randn(N, D_out), requires_grad=False)
```

```
# Construct our model by instantiating the class defined above
model = TwoLayerNet(D_in, H, D_out)
```


```
# Construct our Loss function and an Optimizer. The call to model.parameters()
# in the SGD constructor will contain the learnable parameters of the two
# nn.Linear modules which are members of the model.
criterion = torch.nn.MSELoss(size_average=False)
optimizer = torch.optim.SGD(model.parameters(), lr=1e-4)
for t in range(500):
```

```
    # Forward pass: Compute predicted y by passing x to the model
    y_pred = model(x)
```

```
    # Compute and print loss
    loss = criterion(y_pred, y)
    print(t, loss.data[0])
```

```
    # Zero gradients, perform a backward pass, and update the weights.
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
```

What We Build?



```
# -*- coding: utf-8 -*-
import torch
from torch.autograd import Variable
```

```
class TwoLayerNet(torch.nn.Module):
    def __init__(self, D_in, H, D_out):
        """
        In the constructor we instantiate two nn.Linear modules and assign them as
        member variables.
        """
        super(TwoLayerNet, self).__init__()
        self.linear1 = torch.nn.Linear(D_in, H)
        self.linear2 = torch.nn.Linear(H, D_out)
```

```
    def forward(self, x):
        """
        In the forward function we accept a Variable of input data and we must return
        a Variable of output data. We can use Modules defined in the constructor as
        well as arbitrary operators on Variables.
        """
        h_relu = self.linear1(x).clamp(min=0)
        y_pred = self.linear2(h_relu)
        return y_pred
```

```
# N is batch size; D_in is input dimension;
# H is hidden dimension; D_out is output dimension.
N, D_in, H, D_out = 64, 1000, 100, 10
```

```
# Create random Tensors to hold inputs and outputs, and wrap them in Variables
x = Variable(torch.randn(N, D_in))
y = Variable(torch.randn(N, D_out), requires_grad=False)
```

```
# Construct our model by instantiating the class defined above
model = TwoLayerNet(D_in, H, D_out)
```

```
# Construct our loss function and an Optimizer. The call to model.parameters()
# in the SGD constructor will contain the learnable parameters of the two
# nn.Linear modules which are members of the model.
criterion = torch.nn.MSELoss(size_average=False)
optimizer = torch.optim.SGD(model.parameters(), lr=1e-4)
for t in range(500):
```

```
    # Forward pass: Compute predicted y by passing x to the model
    y_pred = model(x)
```

```
    # Compute and print loss
    loss = criterion(y_pred, y)
    print(t, loss.data[0])
```

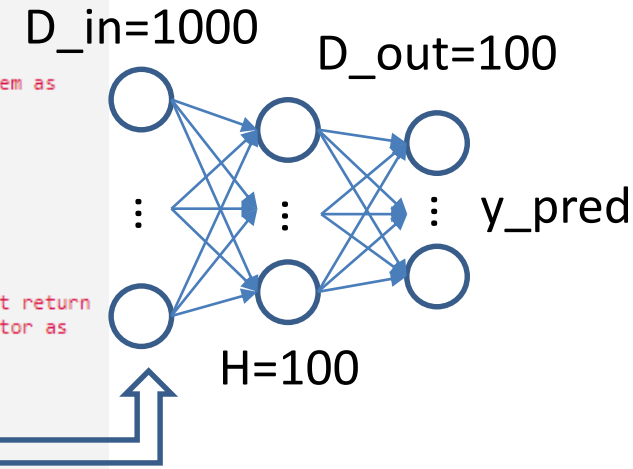
```
    # Zero gradients, perform a backward pass, and update the weights.
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
```

} Don't Update y (y are labels here)

} Construct Our Model

} Optimizer and Loss Function

What We Build?



```
# -*- coding: utf-8 -*-
import torch
from torch.autograd import Variable
```

```
class TwoLayerNet(torch.nn.Module):
    def __init__(self, D_in, H, D_out):
        """
        In the constructor we instantiate two nn.Linear modules and assign them as
        member variables.
        """
        super(TwoLayerNet, self).__init__()
        self.linear1 = torch.nn.Linear(D_in, H)
        self.linear2 = torch.nn.Linear(H, D_out)
```

```
    def forward(self, x):
        """
        In the forward function we accept a Variable of input data and we must return
        a Variable of output data. We can use Modules defined in the constructor as
        well as arbitrary operators on Variables.
        """
        h_relu = self.linear1(x).clamp(min=0)
        y_pred = self.linear2(h_relu)
        return y_pred
```

```
# N is batch size; D_in is input dimension;
# H is hidden dimension; D_out is output dimension.
N, D_in, H, D_out = 64, 1000, 100, 10
```

```
# Create random Tensors to hold inputs and outputs, and wrap them in Variables
x = Variable(torch.randn(N, D_in))
y = Variable(torch.randn(N, D_out), requires_grad=False)
```

```
# Construct our model by instantiating the class defined above
model = TwoLayerNet(D_in, H, D_out)
```

```
# Construct our loss function and an Optimizer. The call to model.parameters()
# in the SGD constructor will contain the learnable parameters of the two
# nn.Linear modules which are members of the model.
criterion = torch.nn.MSELoss(size_average=False)
optimizer = torch.optim.SGD(model.parameters(), lr=1e-4)
for t in range(500):
```

```
    # Forward pass: Compute predicted y by passing x to the model
    y_pred = model(x)
```

```
    # Compute and print loss
    loss = criterion(y_pred, y)
    print(t, loss.data[0])
```

```
    # Zero gradients, perform a backward pass, and update the weights.
    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
```

} Don't Update y (y are labels here)

} Construct Our Model

} Optimizer and Loss Function

Define modules
(must have)

Build network
(must have)

Reset Gradient
Backward
Update Step

Concepts of PyTorch

- Modules of PyTorch

Data:

- Tensor
- Variable (for Gradient)

Function:

- NN Modules
- Optimizer
- Loss Function
- **Multi-Processing**

- Basic Method

- torch.nn.DataParallel
- Recommend by PyTorch

- Advanced Methods

- torch.multiprocessing
- Hogwild (async)

Multi-GPU Processing

- `torch.nn.DataParallel`
 - `gpu_id = '6,7'`
 - `os.environ['CUDA_VISIBLE_DEVICES'] = gpu_id`
 - `net = torch.nn.DataParallel(model, device_ids=[0, 1, 2])`
 - `output = net(input_var)`
- **Important Notes:**
 - Device_ids must start from 0
 - $(\text{batch_size} / \text{GPU_size})$ must be integer


Saving Models

- First Approach (Recommend by PyTorch)
 - `# save only the model parameters`
 - `torch.save(the_model.state_dict(), PATH)`
 - `# load only the model parameters`
 - `the_model = TheModelClass(*args, **kwargs)`
 - `the_model.load_state_dict(torch.load(PATH))`
- Second Approach
 - `torch.save(the_model, PATH) # save the entire model`
 - `the_model = torch.load(PATH) # load the entire model`

Recurrent Neural Network (RNN)

```
class RNN(nn.Module):  
  
    # you can also accept arguments in your model constructor  
    def __init__(self, data_size, hidden_size, output_size):  
        super(RNN, self).__init__()  
  
        self.hidden_size = hidden_size  
        input_size = data_size + hidden_size  
  
        self.i2h = nn.Linear(input_size, hidden_size)  
        self.h2o = nn.Linear(hidden_size, output_size)  
  
    def forward(self, data, last_hidden):  
        input = torch.cat((data, last_hidden), 1)  
        hidden = self.i2h(input)  
        output = self.h2o(hidden)  
        return hidden, output  
  
rnn = RNN(50, 20, 10)
```

```
loss_fn = nn.MSELoss()  
  
batch_size = 10  
TIMESTEPS = 5  
  
# Create some fake data  
batch = Variable(torch.randn(batch_size, 50))  
hidden = Variable(torch.zeros(batch_size, 20))  
target = Variable(torch.zeros(batch_size, 10))  
  
loss = 0  
for t in range(TIMESTEPS):  
    # yes! you can reuse the same network several times,  
    # sum up the losses, and call backward!  
    hidden, output = rnn(batch, hidden)  
    loss += loss_fn(output, target)  
loss.backward()
```




Recurrent Neural Network (RNN)

```
class RNN(nn.Module):  
  
    # you can also accept arguments in your model constructor  
    def __init__(self, data_size, hidden_size, output_size):  
        super(RNN, self).__init__()  
  
        self.hidden_size = hidden_size  
        input_size = data_size + hidden_size  
  
        self.i2h = nn.Linear(input_size, hidden_size)  
        self.h2o = nn.Linear(hidden_size, output_size)  
  
    def forward(self, data, last_hidden):  
        input = torch.cat((data, last_hidden), 1)  
        hidden = self.i2h(input)  
        output = self.h2o(hidden)  
        return hidden, output  
  
rnn = RNN(50, 20, 10)
```

```
loss_fn = nn.MSELoss()  
  
batch_size = 10  
TIMESTEPS = 5  
  
# Create some fake data  
batch = Variable(torch.randn(batch_size, 50))  
hidden = Variable(torch.zeros(batch_size, 20))  
target = Variable(torch.zeros(batch_size, 10))  
  
loss = 0  
for t in range(TIMESTEPS):  
    # yes! you can reuse the same network several times,  
    # sum up the losses, and call backward!  
    hidden, output = rnn(batch, hidden)  
    loss += loss_fn(output, target)  
loss.backward()
```

**Same module (i.e. same parameters)
among the time**



Transfer Learning

- Freeze the parameters of original model
 - `requires_grad = False`
- Then add your own modules

```
model = torchvision.models.resnet18(pretrained=True)
for param in model.parameters():
    param.requires_grad = False
# Replace the last fully-connected layer
# Parameters of newly constructed modules have requires_grad=True by default
model.fc = nn.Linear(512, 100)

# Optimize only the classifier
optimizer = optim.SGD(model.fc.parameters(), lr=1e-2, momentum=0.9)
```

Comparison with TensorFlow

| Properties | TensorFlow | PyTorch |
|------------------------------|-------------------------------------|------------------|
| Graph | Static Dynamic (TensorFlow Fold) | Dynamic |
| Ramp-up Time | - | Win |
| Graph Creation and Debugging | - | Win |
| Feature Coverage | Win | Catch up quickly |
| Documentation | Tie | Tie |
| Serialization | Win (support other lang.) | - |
| Deployment | Win (Cloud & Mobile) | - |
| Data Loading | - | Win |
| Device Management | Win | Need .cuda() |
| Custom Extensions | - | Win |

Summarized from <https://awni.github.io/pytorch-tensorflow/>

Remind: Platform & Final Project

Thank You~!